

V11 Final Validity and Claims Audit

Future Internet submission candidate

Status: PASS. All 56 automated checks passed.

Headline findings

- The compact trace file has 100 public-viewer records, 95 with positive token counts, divided chronologically into 60 search and 35 certification records.
- The locked absolute-SLO study certifies 0 of 120 method runs. Universal abstention is retained as a negative result.
- The separately locked independent-seed noninferiority follow-up certifies 50/60 DR-BGS runs and 33/60 guarded-GP runs.
- Every certified noninferiority run exactly recovers the exhaustive training-selected policy. This is retrospective validation, not a deployable optimality certificate.
- Search plus certification uses 83.95% fewer replications than exhaustive training evaluation plus the same certification batch.

Validity boundary

Arrival times and token counts are trace-derived. Privacy labels, pressure conversion, and moving-block replay are synthetic. The study uses a compact public slice rather than the full BurstGPT corpus. Exact binomial bounds are conditional on independent simulator replays from the declared generator. They are not claims about live production safety. The follow-up was designed after the strict-study result and is not described as external confirmation.

Key numerical checks

Quantity	Audited value
Trace rows / positive rows	100 / 95
Policy surfaces	$12 \times 190 = 2,280$
Strict method runs	120
Absolute certified runs	0
DR-BGS noninferiority certified	50/60
Guarded-GP noninferiority certified	33/60
Search+certification reduction	83.9506%
Python sources compiled	all
Missing or unused citations	none
Overfull boxes	none

Provenance checks

The primary configuration, compact trace CSV, and execution code match the archived SHA-256 values. A portable chronology snapshot preserves the working archive metadata used to record that the primary lock preceded strict result generation and that the follow-up lock preceded its result files. This establishes internal archive chronology and tamper detection; it is not third-party preregistration or trusted timestamping.

Language and claims

Selected promotional and formulaic markers were absent. The manuscript states the failed absolute certificate directly, avoids the term “breakthrough,” and distinguishes deployment certification from retrospective comparison with an exhaustive reference.

Conclusion

The V11 manuscript and its trace-risk results are internally consistent. The scientific package is suitable for external review after the public repository and archival DOI are activated and verified.